

# **System Design – Parallel Programming [05]**

Hochschule  
für Technik  
Stuttgart

Marcus Deiningner  
WS 2021/22

# Overview

- Motivation / Requirements
- Use Cases
- Threads
- Creating Threads
- Controlling Threads
- Master/Worker
- Producer/Consumer
- Pipeline

# Motivation

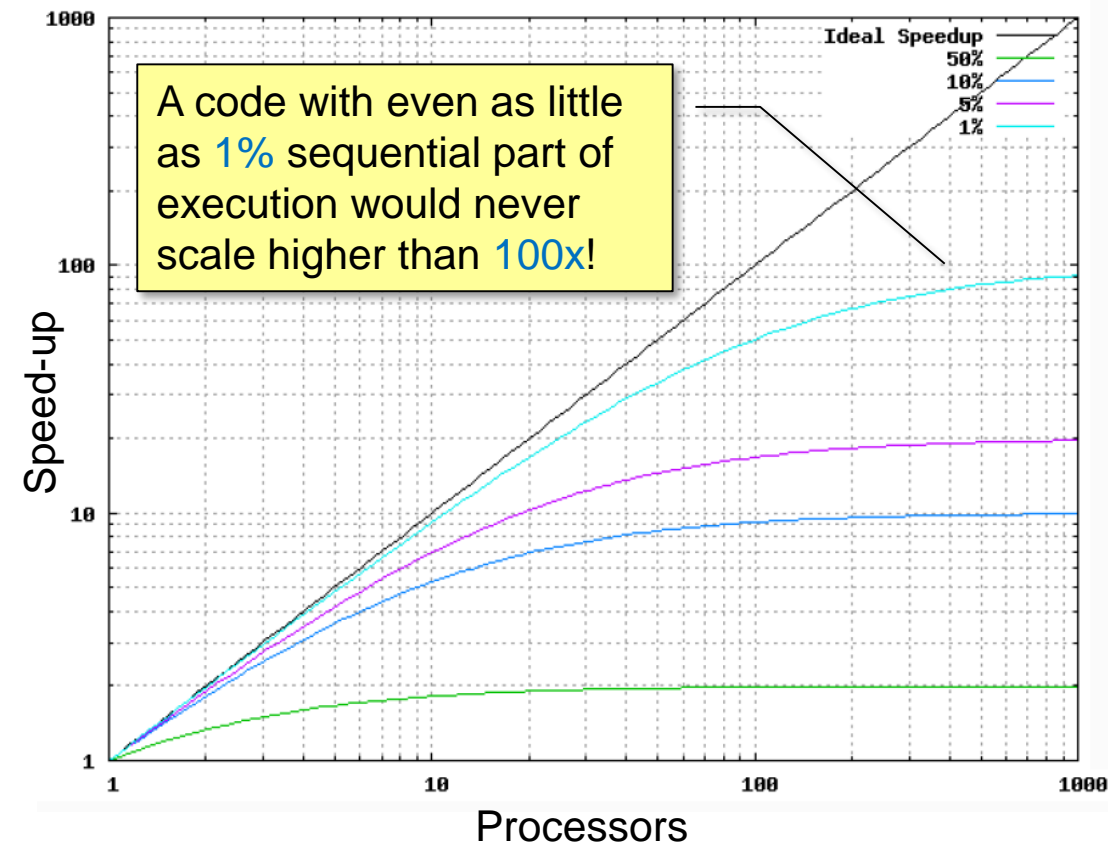
## Parallel Programming

- to run programs in shorter period of time
- to solve with higher accuracy in given time
- to solve a computational problem at all (e.g. because of size of data)

Theoretically restricted by

- Amdahl's Law
- Gustafson's Law

# Amdahl's Law (fixed Problem Size)



- Code has a sequential (s) and a parallel (p) proportion, with  $s + p = 1$

- Total execution Time:


$$T = Ts + Tp$$

- Execution Time with n Processors:  $T_n = Ts + \frac{Tp}{n}$

- Speed-up:

$$\frac{T}{T_n} = \frac{T}{Ts + \frac{Tp}{n}} = \frac{1}{s + \frac{1-s}{n}}$$

# Gustafson's Law (fixed Problem Time)



- Code has a sequential ( $s$ ) and a parallel ( $p$ ) proportion, with  $s + p = 1$
- Ideal execution time running in parallel on  $n$  Processors:  $T_n = Ts + Tp$
- Execution time running this code sequentially:  
 $T = Ts + nTp$
- Speed-up:

$$\frac{T}{T_n} = \frac{Ts + nTp}{Ts + Tp}$$
$$= s + n(s - 1)$$

Figure: [www.researchgate.net](http://www.researchgate.net)

# Which Law is Right?

- Amdahl's Law

- overly pessimistic
- states that the possibilities are very limited

- Gustafson's Law

- overly optimistic
  - states that the possibilities are nearly linear
- the truth is somewhere in the middle

# Requirements for Parallelization

- **Concurrency**
  - multiple tasks can be done at the same time
  - data may be worked on independently (for a longer time)
- **Scalability**
  - if tasks are dependent on each other, there should be little communication and synchronization
- **Locality**
  - if data is exchanged (shared resources /explicit communication), it should be done rarely – i.e. keep processing local.

The worst algorithm will have dependencies in every step, need to exchange all data between all involved processes.

# Processes and Threads

- **Processes**

- separated and isolated tasks with different memory
- heavy-weight, difficult communication

- **Threads**

- independent tasks sharing memory and working on common data
- light-weight, easy communication
- Java natively supports Threads → we will concentrate on **Java-Threads**



# Use Cases

Thread Model	Description
Delegation or Master/Worker	<ul style="list-style-type: none"><li>■ Central <b>Master-Thread</b> creates <b>Worker-Threads</b></li><li>■ Master assigns each worker tasks</li><li>■ Master may wait until workers complete</li></ul>
Producer-Consumer	<ul style="list-style-type: none"><li>■ <b>Producer-Threads</b> produce data</li><li>■ Data is stored in a shared memory</li><li>■ <b>Consumer-Threads</b> consume the data</li></ul>
Pipeline or Pipes & Filters	<ul style="list-style-type: none"><li>■ Processing is done in several stages</li><li>■ Each stage is <b>Filter-Thread</b> that performs work on a unit of input.</li></ul>
Peer-to-Peer	<ul style="list-style-type: none"><li>■ <b>Peer-Threads</b> have equal status and act on their own</li><li>■ working on shared data or private data</li></ul>

# Creating Threads in Java

## Direct creation

- ☑ implement a **Runnable** and start it through a Thread
- create a **Thread**-subclass and start it directly
- implement an anonymous adapter
- ☑ supply a lambda-expression
- implement a named adapter
- **High-level** creation
  - ☑ create an **Executor** and supply a **Runnable**
  - ☑ create an **ExecutorService** and supply a **Callable**
- ☑ *preferred approaches*

# Direct Creation: A “Runnable” Class

Main-program

```
public class Target
    implements Runnable {

    public void run() {
        ... //do something in parallel
    }
}
```

this is the *one* task which should run in parallel

```
Target target = new Target();
Thread myThread = new Thread(target);
myThread.start();
```

forks a new thread and starts “run()”

The Thread-class hosts the task.

# Direct Creation: Lambda Expression

This is the *one* task which should run in parallel – it is **void** and has **no parameters** (just like “run”)

```
public void doSomething() {  
    ... //do something in parallel  
}  
  
...  
Thread myThread = new Thread(() -> doSomething());  
myThread.start();
```

forks a new thread and starts “doSomething()”

The Thread-class hosts the task.

The lambda-Expression

# Typical Structures of run() or doSomething()

```
public void run(){  
    while (true) {  
        try {  
            Thread.sleep(1000);  
        } catch (InterruptedException ignored){}  
        some actions;  
    }  
}
```

an endless loop  
running in the  
Thread after  
start

suspend for  
1000 ms

this is the actual  
task, which  
should be  
performed

# Low-Level-Control of Threads – 1

- start / (stop)
  - `thread.start()` for **starting** a thread
  - ~~`thread.stop()`~~ for stopping – deprecated!
- wait / release
  - `thread.join()` **blocking wait** until the thread has finished
  - `thread.yield()` **releases** the thread

# Low-Level-Control of Threads – 2

- suspend/resume with the help of flags
  - `anObject.wait()` **suspends** an object, waits for `notify()`
  - `anObject.notify()` **resumes** an object
  - `anObject.notifyAll()` **resumes** all objects
  - `Thread.sleep(time)` **suspends** the current thread
  - `thread.interrupt()` sends an **interrupt** to an e.g. sleeping thread

# Low-Level-Control of Threads – 3

## **synchronized** – Method modifier

- a **caller** entering this method **locks** the **whole object**
- all **other callers** trying to access any other **synchronized parts** of this object have to **wait**
- the **lock** is released when the caller **leaves** the method or calls **wait()**
- typically used for **exclusive data access** (“race-conditions”)

## **synchronized** – Block

- more fine-grained

```
synchronized(anObject){ // mostly this
    exclusive actions;
}
```



# Controlling a Thread – 1

```
public class Controlled{  
  
    private Thread thread = new Thread(() -> runInParallel())  
  
    private boolean suspended = false;  
    private long sleepTime = 1000;  
  
    public void runInParallel() {  
        while (true) {  
            try {  
                synchronized(this){  
                    while(suspended)  
                        this.wait();  
                }  
                Thread.sleep(sleepTime);  
            } catch (InterruptedException ignored){}  
            some actions;  
        }  
    }  
    ...  
}
```

an endless loop  
running in parallel  
after start

if the thread wakes  
up from **wait**, he  
first checks his state  
and may go to **wait**  
again (if falsely  
woken up)

this is the actual  
task, which  
should be  
performed

Sleep for a given  
time. This may be  
interrupted.

# Controlling a Thread – 2 (continued)

```
public class Controlled{  
    ...  
    public void setSleepTime(long sleepTime) {  
        thread.interrupt();  
        this.sleepTime = sleepTime;  
    }  
}
```

wake up from  
sleep

```
    public synchronized void start(){  
        if(!thread.isAlive())  
            thread.start();  
        if(suspended){  
            suspended = false;  
            this.notify();  
        }  
    }  
}
```

guarantee that  
this method is  
used by only  
one process

```
    public synchronized void stop(){  
        if(!suspended)  
            suspended = true;  
    }  
}
```

wake up from  
wait

```
    public boolean isRunning(){  
        return thread.isAlive() & !suspended;  
    }  
    ...  
}
```

send to wait

# High-Level-Creation: Executors

- Setting up a **Thread-Pool**

with a fixed number of threads

```
Executor executor = Executors.newFixedThreadPool(3);
```

or system-controlled

```
Executor executor = Executors.newCachedThreadPool();
```

- **Defining and starting** the Threads

```
List<Runnable> runnables = ...
```

```
for(Runnable runnable : runnables)  
    executor.execute(runnable);
```

# High-Level-Creation: Executor-Services

- Setting up a Thread-Pool

```
ExecutorService executor =  
    Executors.newCachedThreadPool();
```

- Defining and starting the Threads

```
Callable<Type> callable = () -> returnSomeType();
```

**Future** is a result available later, when the thread is finished

```
Future<Type> future = executor.submit(callable);
```

- Accessing the result

`future.isDone()` → **true** when the result is available

`Type value = future.get();` blocks until the result is available and retrieves it.

- Stopping the threads

```
executor.shutdown();
```

# High-Level Access Control

- **thread-safe** data-structures
  - AtomicInteger with atomic operations
  - AtomicIntegerArray, AtomicReference<V>, ...
- **thread-safe** collections
  - Collections.*synchronized*...(...) returns the respective thread-safe implementation

```
List<Type> l = new ArrayList<>();  
List<Type> sl = Collections.synchronizedList(l);
```

# Use Case: Master/Worker

- **Divide** the task into (independent) segments
- **Low-Level**
  - Master creates a **Thread** for each segment and starts it
  - Master waits with **join()** for each thread
  - If workers access/change shared data use **synchronized** methods
- **High-Level**
  - Master creates **Executor** or **ExecutorService**
  - assigns **Runnable**s or **Callable**s for each segment
  - retrieves results from **Futures**

# Master/Worker – Sample

```
public class Worker {  
  
    private String name;  
    private int result;  
  
    public Worker(String name) {  
        super();  
        this.name = name;  
    }  
  
    public int getResult() {  
        return result;  
    }  
  
    public void doSomething(int n) {  
        result = 0;  
        while(n > 0) {  
            System.out.println(name);  
            try {  
                Thread.sleep(1000);  
            } catch (InterruptedException e) {  
                e.printStackTrace();  
            }  
            n--;  
            result++;  
        }  
    }  
}
```

```
public class Master {  
  
    public static void main(String[] args)  
        throws InterruptedException {  
        Worker w1 = new Worker("Worker 1");  
        Worker w2 = new Worker("Worker 2");  
  
        Thread t1 = new Thread(() -> w1.doSomething(5));  
        Thread t2 = new Thread(() -> w2.doSomething(10));  
  
        t1.start();  
        t2.start();  
        t1.join();  
        t2.join();  
  
        System.out.println("Result 1: " + w1.getResult());  
        System.out.println("Result 2: " + w2.getResult());  
    }  
}
```


Assign work and start.

Wait for workers to finish.

Get results.

# Use Case: Consumer/Producer

- a producer thread sends data into a queue
- a consumer thread takes the data from a queue
- a producer has to wait if the queue is full
- a consumer has to wait if the queue is empty





# Consumer/Producer – Sample – 1

```
public class Queue {  
  
    public static final int MAXQUEUE = 5;  
    private List<String> messages = new ArrayList<>();  
  
    //called by the Producer  
    public synchronized void putMessage(String message) {  
        while(messages.size() >= MAXQUEUE)  
            try {  
                wait();  
            } catch (InterruptedException e) {}  
        messages.add(message);  
        notifyAll();  
    }  
  
    //called by the Consumer  
    public synchronized String getMessage() {  
        while(messages.size() == 0)  
            try {  
                notifyAll();  
                wait();  
            } catch (InterruptedException e) {}  
        String message = messages.remove(0);  
        notifyAll();  
        return message;  
    }  
}
```

This is the  
shared data: it  
is synchronized,  
not “threaded”

secured data  
access

resume a  
possible waiting  
producer

wait for a queue  
to be filled

resume possible  
producers, as there  
is now space in the  
queue

# Consumer/Producer – Sample – 2

```
public class NamedProducer implements Runnable {  
  
    private String name;  
    private Queue queue;  
  
    public NamedProducer(String name, Queue queue) {  
        super();  
        this.name = name;  
        this.queue = queue;  
    }  
  
    @Override  
    public void run() {  
        while(true){  
            queue.putMessage(new Date().toString());  
            try {  
                Thread.sleep(1000);  
            } catch (InterruptedException e) {}  
        }  
    }  
}
```

This is to make the **producers / consumers** run in parallel later on – otherwise there would be no sense in synchronizing the data.

The **producer** creates new products and puts it into the queue.

This is just for demo effect – to force some actual delay

# Consumer/Producer – Sample – 3

```
public class NamedConsumer implements Runnable {  
  
    private String name;  
    private Queue queue;  
  
    public NamedConsumer(String name, Queue queue) {  
        super();  
        this.name = name;  
        this.queue = queue;  
    }  
  
    @Override  
    public void run() {  
        while(true){  
            String message = queue.getMessage();  
            System.out.println(name + " got message; " + message);  
            try {  
                Thread.sleep(2000);  
            } catch (InterruptedException e) {}  
        }  
    }  
}
```

The consumer receives products from the queue.

This is just for demo effect – to force some actual delay

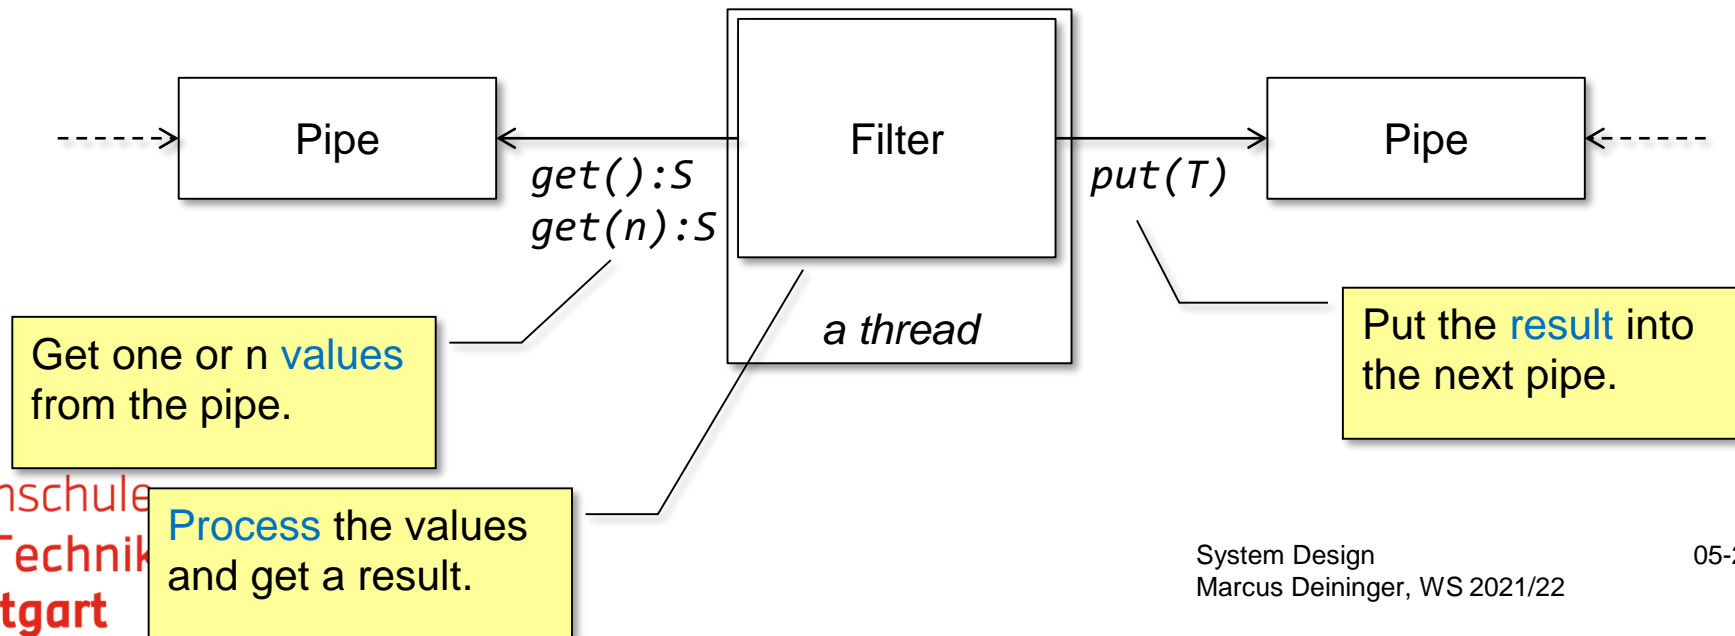
# Controlling Data Access – 4

Main program for  
starting the  
threads.

```
public class Main {  
  
    public static void main(String[] args){  
        Queue queue = new Queue();  
  
        NamedProducer producer1 = new NamedProducer("P1", queue);  
        new Thread(producer1).start();  
        NamedProducer producer2 = new NamedProducer("P2", queue);  
        new Thread(producer2).start();  
        NamedConsumer consumer1 = new NamedConsumer("C1", queue);  
        new Thread(consumer1).start();  
        NamedConsumer consumer2 = new NamedConsumer("C2", queue);  
        new Thread(consumer2).start();  
    }  
}
```

# Use Case: Pipeline

- A **Pipe** is a **Queue** with synchronized access
- A **Filter** is a common **Consumer/Producer** with
  - its own thread
  - an assigned input queue from which he consumes
  - an assigned output queue to which he produces



# Pipeline – Sample (→ Exercise)

Converting Data  
from Source S to  
Target T

```
public class BasicPipe<T> implements Pipe<T>{

    private static final int MAXQUEUE = 5;
    private List<T> queue = new LinkedList<>();

    //called by the Producer
    @Override
    public synchronized void put(T t) {...}

    //called by the Consumer
    @Override
    public synchronized T get() {...}

    //called by the Consumer
    @Override
    public synchronized List<T> get(int n) {...}

    //called by the Consumer
    @Override
    public synchronized List<T> getAll(int n) {...}

    public synchronized List<T> getAll() {...}

}
```

```
public abstract class Filter<S, T> implements Runnable {

    protected Pipe<S> in = null;
    protected Pipe<T> out = null;
    protected boolean running = true;

    public boolean isRunning() {
        return running;
    }

    public void setIn(Pipe<S> in) {
        this.in = in;
    }

    public void setOut(Pipe<T> out) {
        this.out = out;
    }

    public abstract void process();

    @Override
    public void run() {
        while(isRunning())
            process();
    }

}
```

The actual action  
on the data work-  
ing on in and out

The Queue of  
Producer /  
Consumer

The Producer  
*and* Consumer

# Pipeline – Sample (→ Exercise)

```
public class Pipeline {  
  
    private List<Thread> threads;  
  
    private Pipeline(List<Thread> threads) {  
        this.threads = threads;  
    }  
  
    public static class Segment<S>{  
  
        private List<Thread> threads;  
        private Pipe<S> in = null;  
  
        ...  
  
        public <T> Segment<T> filter(Filter<S, T> filter) {  
            filter.in = in;  
            Pipe<T> p = new BasicPipe<T>();  
            filter.out = p;  
            threads.add(new Thread(filter));  
            Segment<T> next = new Segment<>(p, threads);  
            return next;  
        }  
        ...  
    }  
  
    public void start() {  
        for(Thread thread : threads)  
            thread.start();  
    }  
}
```

The list of  
Threads hosting  
the filters

One pipeline-  
segment

Connect to  
previous pipe.

Create following  
pipe and connect

Host filter

Next segment  
with new  
“dangling” pipe

Start all threads.

# Pipeline – Sample (→ Exercise)

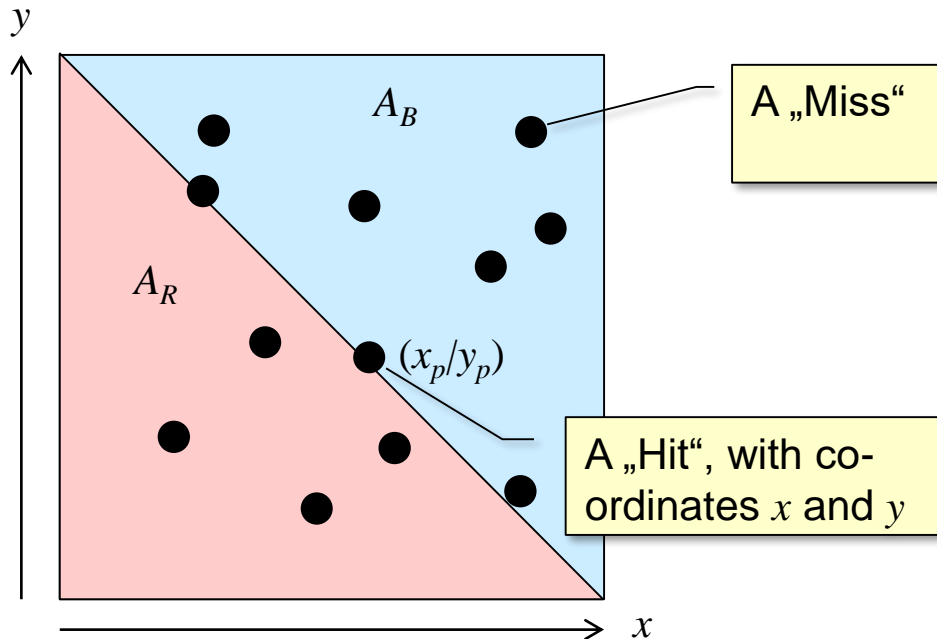
```
public class Sample {  
  
    public static String keep(String s) {  
        return s;  
    }  
  
    public static void main(String[] args) {  
        Pipeline p1 = Pipeline.source(new TextFileSource("test.txt"))  
                                .map(s -> keep(s))  
                                .drain(new OutputDrain<String>());  
  
        p1.start();  
    }  
}
```



# Use Case: Peer-to-Peer

- mixture of the previous use-cases
- not really discussed here

# Geometrical Probability



Chances are  
50%, that a  
random shot ends  
up in  $A_R$